

wrangle_act

October 23, 2019

1 Gathering Data

```
In [0]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive")

```
In [0]: import pandas as pd
import requests
import tweepy
import json
import sys
import datetime as dt
import matplotlib.pyplot as plt
import re
import numpy as np
```

```
In [0]: fn_twitter_archive_enhanced = '/content/drive/My Drive/Colab Notebooks/Data Wrangling/Pr
fn_img_pred_local_file_name = '/content/drive/My Drive/Colab Notebooks/Data Wrangling/Pr
fn_local_json_file = '/content/drive/My Drive/Colab Notebooks/Data Wrangling/Project/Da
fn_error_log = '/content/drive/My Drive/Colab Notebooks/Data Wrangling/Project/Data/json
fn_local_json_file1 = '/content/drive/My Drive/Colab Notebooks/Data Wrangling/Project/Da
fn_final_clean_file = '/content/drive/My Drive/Colab Notebooks/Data Wrangling/Project/Da
```

```
img_pred_url = 'https://d17h27t6h515a5.cloudfront.net/topher/2017/August/599fd2ad_image-
```

```
consumer_key = 'MY CONSUMER KEY'
consumer_secret = 'MY CONSUMER SECRET KEY'
access_token = 'MY ACCESS TOKEN'
access_secret = 'MY ACCESS SECRET'
```

```
In [0]: # Load the twitter enhance archive data
df_twitter_archive = pd.read_csv(fn_twitter_archive_enhanced)
df_twitter_archive.head()
```

```
Out[0]:
```

	tweet_id	in_reply_to_status_id	...	pupper	puppo
0	892420643555336193	NaN	...	None	None

1	892177421306343426	NaN	...	None	None
2	891815181378084864	NaN	...	None	None
3	891689557279858688	NaN	...	None	None
4	891327558926688256	NaN	...	None	None

[5 rows x 17 columns]

```
In [0]: # Download the image prediction data
response = requests.get(img_pred_url)
```

```
In [0]: # Store the data in a local file
with open(fn_img_pred_local_file_name, mode='wb') as file:
    file.write(response.content)
```

```
In [0]: # Load the data into a dataframe
df_image_pred = pd.read_csv(fn_img_pred_local_file_name, sep='\t')
df_image_pred.head()
```

```
Out[0]:
```

	tweet_id	...	p3_dog
0	666020888022790149	...	True
1	666029285002620928	...	True
2	666033412701032449	...	True
3	666044226329800704	...	True
4	666049248165822465	...	True

[5 rows x 12 columns]

1.0.1 Code for Downloading Twitter data. This did not work. I had Rate limit issues

```
In [0]: # Setup to access data from twitter
auth = tweepy.OAuthHandler(consumer_key, consumer_secret)
auth.set_access_token(access_token, access_secret)

api = tweepy.API(auth, wait_on_rate_limit=True, wait_on_rate_limit_notify=True)

# For each tweet, get the tweet details
logf = open(fn_error_log, 'w')
outfile = open(fn_local_json_file, 'w')
for index, row in df_twitter_archive.iterrows():
    t_id = row['tweet_id']
    try:
        data = api.get_status(t_id, tweet_mode='extended')._json + '\n'
        json.dump(data, outfile)
        logf.write('Success|' + str(t_id) + '\n')
        print(str(t_id))
    except:
        logf.write('Error|' + str(t_id) + '\n')

logf.close()
```

```

outfile.close()
print('Done')

```

1.0.2 Alternate to twitter download. Used the file provided

```

In [0]: # Gather only those tweets which are not re-tweets
        # if line contains 'retweeted_status' key, it would mean its a retweet
        j_data = []
        with open(fn_local_json_file1, 'r') as j_file:
            for line in j_file:
                if 'retweeted_status' not in line:
                    j_content = json.loads(line)
                    j_data.append([j_content['id'], j_content['retweet_count'], j_content['favorite_count']])

df_jdata = pd.DataFrame(j_data, columns=['tweet_id', 'retweet_count', 'favorite_count'])
df_jdata.head()

```

```

Out[0]:
      tweet_id  retweet_count  favorite_count
0  89242064355336193         8853         39467
1  892177421306343426         6514         33819
2  891815181378084864         4328         25461
3  891689557279858688         8964         42908
4  891327558926688256         9774         41048

```

2 Assessing Data - eight (8) quality issues and two (2) tidiness issues

```

In [0]: # Create a copy of all the extracted/downloaded data
        df_twitter_arch_clean = df_twitter_archive.copy()
        df_image_pred_clean = df_image_pred.copy()
        df_jdata_clean = df_jdata.copy()

```

```

In [0]: df_twitter_arch_clean.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2356 entries, 0 to 2355
Data columns (total 17 columns):
tweet_id                2356 non-null int64
in_reply_to_status_id   78 non-null float64
in_reply_to_user_id     78 non-null float64
timestamp               2356 non-null object
source                 2356 non-null object
text                   2356 non-null object
retweeted_status_id     181 non-null float64
retweeted_status_user_id 181 non-null float64
retweeted_status_timestamp 181 non-null object
expanded_urls           2297 non-null object
rating_numerator        2356 non-null int64

```

```

rating_denominator      2356 non-null int64
name                    2356 non-null object
doggo                   2356 non-null object
floofer                 2356 non-null object
pupper                  2356 non-null object
puppo                   2356 non-null object
dtypes: float64(4), int64(3), object(10)
memory usage: 313.0+ KB

```

```
In [0]: df_twitter_arch_clean.head(5)
```

```

Out[0]:
      tweet_id  in_reply_to_status_id  ... pupper puppo
0  892420643555336193              NaN  ...   None   None
1  892177421306343426              NaN  ...   None   None
2  891815181378084864              NaN  ...   None   None
3  891689557279858688              NaN  ...   None   None
4  891327558926688256              NaN  ...   None   None

[5 rows x 17 columns]

```

```
In [0]: # check if any of these tweets are re-tweeted. These need to be removed
df_twitter_arch_clean.retweeted_status_id.notnull().sum()
```

```
Out[0]: 181
```

```

In [0]: # On visual checks it was noticed that there were multiple tweets where the picture was
# It had the text, 'we only rate dogs' lets check how many do we have like this
str_search = 'WE ONLY RATE DOGS'
df_twitter_arch_clean['text_uppercase'] = df_twitter_arch_clean.text.str.upper()
df_twitter_arch_clean[df_twitter_arch_clean['text_uppercase'].str.contains(str_search)].

```

```
Out[0]: 63
```

```

In [0]: # Drop the additional column
df_twitter_arch_clean.drop('text_uppercase', axis=1, inplace=True)

```

```

In [0]: # Rating numerator and denominator do not have null values
# Are there any rows where the denominator is not 10
df_temp = df_twitter_arch_clean[df_twitter_arch_clean['rating_denominator']!=10]
df_temp.head()

```

```

Out[0]:
      tweet_id  in_reply_to_status_id  ... pupper puppo
313  835246439529840640      8.352460e+17  ...   None   None
342  832088576586297345      8.320875e+17  ...   None   None
433  820690176645140481              NaN  ...   None   None
516  810984652412424192              NaN  ...   None   None
784  775096608509886464              NaN  ...   None   None

[5 rows x 17 columns]

```

```
In [0]: # Is this because the data is incorrect or due to the fact the data was extracted incorr
df_temp[['tweet_id', 'text', 'rating_denominator', 'rating_numerator']].head(5)
```

```
Out[0]:
```

	tweet_id	...	rating_numerator
313	835246439529840640	...	960
342	832088576586297345	...	11
433	820690176645140481	...	84
516	810984652412424192	...	24
784	775096608509886464	...	9

[5 rows x 4 columns]

Visual checks show that the data is not being picked up corectly from the text field. The field rating_denominator is incorrectly picked up. Needs to be fixed. This will impact rating_numerator as well.

```
In [0]: # In case we are picking up the value from the text, lets check if there are any decimal
df_twitter_arch_clean[df_twitter_arch_clean.text.str.contains(r"(\d+\.\d*/\d+)")][['tex
```

```
/usr/local/lib/python3.6/dist-packages/ipykernel_launcher.py:1: UserWarning: This pattern has ma
"""Entry point for launching an IPython kernel.
```

```
Out[0]:
```

	text	...	rating_denominator
45	This is Bella. She hopes her smile made you sm...	...	10
340	RT @dog_rates: This is Logan, the Chow who liv...	...	10
695	This is Logan, the Chow who lived. He solemnly...	...	10
763	This is Sophie. She's a Jubilant Bush Pupper.	10
1689	I've been told there's a slight possibility he...	...	10
1712	Here we have uncovered an entire battalion of	10

[6 rows x 3 columns]

On visually checking there are decimal values. The numerator column should be a float type

```
In [0]: # Check for the values in rating_numerator as well
df_twitter_arch_clean.rating_numerator.value_counts()
```

```
Out[0]:
```

12	558
11	464
10	461
13	351
9	158
8	102
7	55
14	54
5	37
6	32
3	19

```

4          17
1           9
2           9
420         2
0           2
15          2
75          2
80          1
20          1
24          1
26          1
44          1
50          1
60          1
165         1
84          1
88          1
144         1
182         1
143         1
666         1
960         1
1776        1
17          1
27          1
45          1
99          1
121         1
204         1
Name: rating_numerator, dtype: int64

```

There are values like 420, 0, 165 etc. lets check the actual text values for these tweets

```

In [0]: # Check for tweet texts where the numerator is 0 or greater than 20
df_temp = df_twitter_arch_clean[(df_twitter_arch_clean['rating_numerator']==0) | (df_twitter_arch_clean['rating_numerator']>20)]
df_temp[['tweet_id', 'text', 'rating_numerator']].head(5)

```

```

Out[0]:
   tweet_id  ... rating_numerator
188  855862651834028034  ...          420
189  855860136149123072  ...          666
290  838150277551247360  ...          182
313  835246439529840640  ...          960
315  835152434251116546  ...           0

[5 rows x 3 columns]

```

On checking these values visually as well, seems like these are valid values. No change required.

```
In [0]: # check for any duplicate rows in the dataframe
df_twitter_arch_clean.duplicated().sum()
```

```
Out[0]: 0
```

```
In [0]: # check for the vales in the dog stage
df_twitter_arch_clean.doggo.unique(), df_twitter_arch_clean.floofer.unique(), df_twitter
```

```
Out[0]: (array(['None', 'doggo'], dtype=object),
array(['None', 'floofer'], dtype=object),
array(['None', 'pupper'], dtype=object),
array(['None', 'puppo'], dtype=object))
```

```
In [0]: # check to see if they are mutually unique
df_twitter_arch_clean['combined_puppy_stage'] = (df_twitter_arch_clean.doggo + df_twitter
df_twitter_arch_clean['combined_puppy_stage'] = df_twitter_arch_clean['combined_puppy_st
df_twitter_arch_clean.combined_puppy_stage.value_counts()
```

```
Out[0]:
```

	1976
pupper	245
doggo	83
puppo	29
doggopupper	12
floofer	9
doggopuppo	1
doggofloofer	1

Name: combined_puppy_stage, dtype: int64

```
In [0]: # Check the dog name column
df_twitter_arch_clean.name.value_counts()
```

```
Out[0]:
```

None	745
a	55
Charlie	12
Lucy	11
Oliver	11
Cooper	11
Penny	10
Lola	10
Tucker	10
Winston	9
Bo	9
Sadie	8
the	8
Toby	7
Daisy	7
Bailey	7
an	7
Buddy	7

Jax	6
Dave	6
Jack	6
Scout	6
Leo	6
Stanley	6
Rusty	6
Koda	6
Oscar	6
Bella	6
Milo	6
Chester	5

...

Jaycob	1
Biden	1
Vixen	1
Sandra	1
Jimbo	1
Oreo	1
Katie	1
Taz	1
Mo	1
Zuzu	1
Berkeley	1
space	1
Lillie	1
Venti	1
Zoe	1
Bayley	1
Keet	1
Willem	1
Ivar	1
Dietrich	1
Yukon	1
Patch	1
Laela	1
Wishes	1
Mutt	1
Callie	1
Tove	1
Kobe	1
Grady	1
Jebberson	1

Name: name, Length: 957, dtype: int64

Are names like all, None, an valid values - lets do a visual check. On visually checking, None - there does not seem to be a name in these tweets, seems okay. Should be set to Blank a, all, the, actually - again here it seems to be incorrectly extracted.


```
In [0]: # It seems that the dog names starting with a lower case are incorrect. Lets see some of
df_twitter_arch_clean[df_twitter_arch_clean['name'].str[0].str.islower()]['name'].value_
```

```
Out[0]: a          55
the          8
an           7
very         5
quite        4
one          4
just         4
not          2
actually     2
mad          2
getting      2
his          1
infuriating  1
unacceptable 1
officially   1
space        1
old          1
incredibly   1
life         1
light        1
my           1
such         1
by           1
all          1
this         1
Name: name, dtype: int64
```

```
In [0]: # Lets check the Image prediction dataframe
df_image_pred_clean.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2075 entries, 0 to 2074
Data columns (total 12 columns):
tweet_id      2075 non-null int64
jpg_url       2075 non-null object
img_num       2075 non-null int64
p1            2075 non-null object
p1_conf       2075 non-null float64
p1_dog        2075 non-null bool
p2            2075 non-null object
p2_conf       2075 non-null float64
p2_dog        2075 non-null bool
p3            2075 non-null object
p3_conf       2075 non-null float64
p3_dog        2075 non-null bool
```

```
dtypes: bool(3), float64(3), int64(2), object(4)
memory usage: 152.1+ KB
```

```
In [0]: #check for any duplicates, since there are absolutely no null entries
        df_image_pred_clean.duplicated().sum(), df_image_pred_clean.tweet_id.duplicated().sum()
```

```
Out[0]: (0, 0)
```

```
In [0]: df_image_pred_clean.head()
```

```
Out[0]:
```

	tweet_id	...	p3_dog
0	666020888022790149	...	True
1	666029285002620928	...	True
2	666033412701032449	...	True
3	666044226329800704	...	True
4	666049248165822465	...	True

```
[5 rows x 12 columns]
```

```
In [0]: df_image_pred_clean.p1.value_counts()
```

```
Out[0]:
```

golden_retriever	150
Labrador_retriever	100
Pembroke	89
Chihuahua	83
pug	57
chow	44
Samoyed	43
toy_poodle	39
Pomeranian	38
cocker_spaniel	30
malamute	30
French_bulldog	26
miniature_pinscher	23
Chesapeake_Bay_retriever	23
seat_belt	22
Siberian_husky	20
German_shepherd	20
Staffordshire_bullterrier	20
Cardigan	19
web_site	19
teddy	18
Maltese_dog	18
Shetland_sheepdog	18
Eskimo_dog	18
beagle	18
Lakeland_terrier	17
Rottweiler	17

Shih-Tzu	17
Italian_greyhound	16
kuvasz	16
...	
bow	1
basketball	1
teapot	1
water_buffalo	1
hay	1
cuirass	1
fountain	1
ping-pong_ball	1
ice_lolly	1
grille	1
military_uniform	1
agama	1
desktop_computer	1
African_hunting_dog	1
convertible	1
piggy_bank	1
king_penguin	1
bonnet	1
dining_table	1
nail	1
hare	1
pitcher	1
carousel	1
maillot	1
espresso	1
candle	1
rain_barrel	1
ibex	1
Scotch_terrier	1
trombone	1

Name: p1, Length: 378, dtype: int64

```
In [0]: df_image_pred_clean.tail()
```

```
Out[0]:
```

	tweet_id	...	p3_dog
2070	891327558926688256	...	True
2071	891689557279858688	...	False
2072	891815181378084864	...	True
2073	892177421306343426	...	True
2074	892420643555336193	...	False

[5 rows x 12 columns]

```
In [0]: # Json data file
df_jdata_clean.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2175 entries, 0 to 2174
Data columns (total 3 columns):
tweet_id          2175 non-null int64
retweet_count     2175 non-null int64
favorite_count    2175 non-null int64
dtypes: int64(3)
memory usage: 51.1 KB

```

3 Data Quality and Tidiness Issues

- DataFrame for Archived Twitter data

- 1) Quality Issue - Base on the column retweeted_status_id, we noticed that there are around 181 retweets in this dataset. This should be removed.
- 2) Quality Issue - There were tweets with the Text 'We only rate dogs'. These tweets should be excluded from our analysis.
- 3) Quality Issue - There are rows where the rating_denominator is not 10. On checking visually for these rows, it seems that they were not parsed correctly from the text. Re-extract the Rating denominator from the text
- 4) Quality Issue - For the rows where rating_denominator is incorrect, it was also noticed that the rating_numerator is incorrect. This needs to be correctly extracted.
- 5) Tidiness Issue - Currently the rating_numerator and rating_denominator columns are type objects which would be string. We also noticed that some numerator values could be decimals. Change numerator to float and denominator to int
- 6) Tidiness Issue - The puppy stage columns could have been consolidated to have a categorical value. However, there are 2 rows which have both values. For these 2, concatenate using ','.
- 7) Quality Issue - The puppy stage columns have text like 'None' Can be fixed as Blank.
- 8) Quality Issue - The timestamp column has additional details like +0000 which does not look necessary/correct.
- 9) Tidiness Issue - The timestamp, retweeted_status_timestamp columns are object instead of datetime. Change timestamp to datetime type
- 10) Tidiness Issue - in_reply_to_status_id, in_reply_to_user_id expanded_urls contains null values. These can be excluded from analysis.
- 11) Quality Issue - The dog name column has 'None'. Should be Blank instead
- 12) Quality Issue - Are names like all, None, an valid values. On visually checking, For values where it starts with a lowercase letter: a, all, the, actually - again here it seems to be incorrectly extracted. However on visually checking these records, it seems that there are no names in these tweets. Reset these values to Blank

- DataFrame for Image Prediction data

- 13) Quality Issue - There are 2356 entries in the Twitter Archived dataframe where as here there are only 2075 entries. We are missing approximately around 300 records at the least not considering any mismatches. Since these are missing, they cannot be merged back at the moment.
 - 14) Tidiness Issue - The predicted p1, p2 and p3 have values which are not dogs. Relevant data would be to pick the prediction which is actually a dog with the highest probability. Create a scoring system to ascertain the best probability of a dog
- DataFrame for Json data
- 15) This is the filtered list of tweets which are not re-tweets. This will be the source list for the rest of the data when combined together
 - 16) Tidiness Issue - Pull in the data from Image Prediction and Json data to the twitter archive dataframe for analysis

4 Cleaning the data

```
In [0]: # DEFINE: Issue 1: Quality Issue - Base on the column retweeted_status_id, we noticed th
# CLEAN: Issue 1:
df_twitter_arch_clean = df_twitter_arch_clean[df_twitter_arch_clean.retweeted_status_id.
```

```
In [0]: # TEST CLEANED: Issue 1
df_twitter_arch_clean.retweeted_status_id.notnull().sum()
```

```
Out[0]: 0
```

```
In [0]: # DEFINE: Issue 2: Quality Issue - On visual checks it was noticed that there were multi
# It had the text, 'we only rate dogs'. Remove these rows
# CLEAN: Issue 2:
str_search = 'WE ONLY RATE DOGS'
df_twitter_arch_clean['text_uppercase'] = df_twitter_arch_clean.text.str.upper()
df_twitter_arch_clean = df_twitter_arch_clean[~df_twitter_arch_clean['text_uppercase'].s
```

```
In [0]: # TEST CLEANED: Issue 2
df_twitter_arch_clean[df_twitter_arch_clean['text_uppercase'].str.contains(str_search)].
```

```
Out[0]: 0
```

```
In [0]: # DEFINE:

# Issue 3: Quality Issue - There are rows where the rating_denominator is not 10. On che
# Re-extract the Rating denominator from the text

# Issue 4: Quality Issue - For the rows where rating_denominator is incorrect, it was al

# Issue 5: Quality/Tidiness Issue - Currently the rating_numerator and rating_denominator
# We also noticed that some numerator values could be decimals. Add clean colu
```

```

# CLEAN: Issue 3, 4, 5:
# For each row, get the twitter text, Split by pattern %d/%d, take the one which has 10
df_twitter_arch_clean['numerator_clean'] = 0.0
df_twitter_arch_clean['denominator_clean'] = 0.0
df_twitter_arch_clean.numerator_clean.astype(float)
df_twitter_arch_clean.denominator_clean.astype(float)

for index, row in df_twitter_arch_clean.iterrows():
    s_split = re.findall('\d+\.\d*/\d+', row.text) # this pattern should give us both dec
    for s in s_split:
        num = float(s.split('/')[0])
        denom = float(s.split('/')[1])
        if(denom==10):
            df_twitter_arch_clean.loc[index, 'numerator_clean'] = num
            df_twitter_arch_clean.loc[index, 'denominator_clean'] = denom
            break

```

```

In [0]: # TEST CLEANED:
# Issue 3 and 4
df_twitter_arch_clean[df_twitter_arch_clean['rating_denominator']!=10][['tweet_id', 'text']]

```

```

Out[0]:
      tweet_id  ... denominator_clean
313  835246439529840640  ...           10.0
342  832088576586297345  ...            0.0
433  820690176645140481  ...            0.0
516  810984652412424192  ...            0.0
902  758467244762497024  ...            0.0

```

[5 rows x 6 columns]

```

In [0]: # Check why there are 0 value denominators
df_twitter_arch_clean[df_twitter_arch_clean.denominator_clean==0][['tweet_id', 'text']]

```

```

Out[0]:
      tweet_id  ... text
342  832088576586297345  @docmisterio account started on 11/15/15
433  820690176645140481  The floofs have been released I repeat the flo...
516  810984652412424192  Meet Sam. She smiles 24/7 & secretly aspir...
902  758467244762497024  Why does this never happen at my front door...
1120 731156023742988288  Say hello to this unbelievably well behaved sq...
1228 713900603437621249  Happy Saturday here's 9 puppies on a bench. 99...
1254 710658690886586372  Here's a brigade of puppies. All look very pre...
1274 709198395643068416  From left to right:\nCletus, Jerome, Alejandro...
1351 704054845121142784  Here is a whole flock of puppies. 60/50 I'll ...
1433 697463031882764288  Happy Wednesday here's a bucket of pups. 44/40...
1598 686035780142297088  Yes I do realize a rating of 4/20 would've bee...
1634 684225744407494656  Two sneaky puppies were not initially seen, mo...
1635 684222868335505415  Someone help the girl is being mugged. Several...
1663 682808988178739200  I'm aware that I could've said 20/16, but here...

```

```

1779 677716515794329600 IT'S PUPPERGEDDON. Total of 144/120 ...I think...
1843 675853064436391936 Here we have an entire platoon of puppers. Tot...

```

For the ones that are blank in our extract, on doing a visual check, we noticed some them indeed had valid values. However, some of them did not. For eg. 204/170 is a valid rating. However, text 'account started on 11/15/15' does not give us the right information. Some of them had 50/50 and 13/10 in the same text and the other had 45/50. There was no pattern to distinguish between this. For now, we are considering only ratings with a denominator of 10

```
In [0]: # TEST CLEANED:
```

```

# Issue 5 - Check the column types added
df_twitter_arch_clean.info()

```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2121 entries, 0 to 2355
Data columns (total 21 columns):
tweet_id                2121 non-null int64
in_reply_to_status_id   77 non-null float64
in_reply_to_user_id     77 non-null float64
timestamp               2121 non-null object
source                  2121 non-null object
text                    2121 non-null object
retweeted_status_id     0 non-null float64
retweeted_status_user_id 0 non-null float64
retweeted_status_timestamp 0 non-null object
expanded_urls           2063 non-null object
rating_numerator        2121 non-null int64
rating_denominator      2121 non-null int64
name                    2121 non-null object
doggo                   2121 non-null object
floofer                 2121 non-null object
pupper                  2121 non-null object
puppo                   2121 non-null object
combined_puppy_stage    2121 non-null object
text_uppercase          2121 non-null object
numerator_clean         2121 non-null float64
denominator_clean       2121 non-null float64
dtypes: float64(6), int64(3), object(12)
memory usage: 444.5+ KB

```

```
In [0]: #DEFINE
```

```
# Issue 6: Tidiness Issue - The puppy stage columns could have been consolidated to have
```

```
# Issue 7: Quality Issue - The puppy stage columns have text like 'None' Can be fixed as
```

```
# CLEAN Issue 7:
```

```
#The puppy stage columns have text like 'None' Can be fixed as Blank.
```

```
df_twitter_arch_clean.doggo.replace('None', '', inplace=True)
df_twitter_arch_clean.floofer.replace('None', '', inplace=True)
df_twitter_arch_clean.pupper.replace('None', '', inplace=True)
df_twitter_arch_clean.puppo.replace('None', '', inplace=True)
```

In [0]: # TEST Issue 7

```
df_twitter_arch_clean.doggo.str.contains('None').sum(),df_twitter_arch_clean.floofer.str
```

Out[0]: (0, 0, 0, 0)

In [0]: # CLEAN Issue 6

The puppy stage columns could have been consolidated to have a categorical value. However

```
df_twitter_arch_clean['combined_puppy_stage'] = ''
```

```
for index, row in df_twitter_arch_clean.iterrows():
```

```
    doggo, floofer, pupper, puppo = row.doggo, row.floofer, row.pupper, row.puppo
```

```
    str_dogstage = ''
```

```
    if(doggo!=''):
```

```
        str_dogstage=doggo
```

```
    if(floofer!=''):
```

```
        if(str_dogstage==''):
```

```
            str_dogstage = floofer
```

```
        else:
```

```
            str_dogstage = str_dogstage + ',' + floofer
```

```
    if(pupper!=''):
```

```
        if(str_dogstage==''):
```

```
            str_dogstage = pupper
```

```
        else:
```

```
            str_dogstage = str_dogstage + ',' + pupper
```

```
    if(puppo!=''):
```

```
        if(str_dogstage==''):
```

```
            str_dogstage = puppo
```

```
        else:
```

```
            str_dogstage = str_dogstage + ',' + puppo
```

```
df_twitter_arch_clean.loc[index,'combined_puppy_stage']=str_dogstage
```

In [0]: # TEST Issue 6

Check the value counts

```
df_twitter_arch_clean['combined_puppy_stage'].value_counts()
```

Out[0]:

```
1777
```

```
pupper      224
```

```
doggo        75
```

```
puppo        24
```

```
doggo,pupper  10
```

```
floofer        9
```

```
doggo,puppo    1
```

```
doggo,floofer  1
```

```
Name: combined_puppy_stage, dtype: int64
```



```

In [0]: # DEFINE
        # Issue 8 Quality Issue - The timestamp column has additional details like +0000 which a

        # Issue 9 Tidiness Issue - The timestamp, retweeted_status_timestamp columns are object

        #CLEAN Issues 8 and 9
df_twitter_arch_clean['timestamp_clean'] = pd.to_datetime(df_twitter_arch_clean['timesta
df_twitter_arch_clean['retweeted_status_timestamp_clean'] = pd.to_datetime(df_twitter_ar

In [0]: # TEST Issues 8 and 9
print(str(df_twitter_arch_clean['timestamp'].dtype) + ' ' + str(df_twitter_arch_clean['t
print(str(df_twitter_arch_clean['retweeted_status_timestamp'].dtype) + ' ' + str(df_twit

object datetime64[ns, UTC]
object datetime64[ns]

In [0]: df_twitter_arch_clean.timestamp_clean.head()

Out[0]: 0    2017-08-01 16:23:56+00:00
        1    2017-08-01 00:17:27+00:00
        2    2017-07-31 00:18:03+00:00
        3    2017-07-30 15:58:51+00:00
        4    2017-07-29 16:00:24+00:00
        Name: timestamp_clean, dtype: datetime64[ns, UTC]

In [0]: # DEFINE Issue 12 Quality Issue - Are names like all, None, an valid values. On visually
        # a, all, the, actually - again here it seems to be incorrectly extracted. However on vi
        # Reset these values to Blank

        # CLEAN Issue 12
df_twitter_arch_clean['name_clean'] = df_twitter_arch_clean['name']
for index, row in df_twitter_arch_clean.iterrows():
    str_name_clean = row['name_clean']
    if(str(str_name_clean)[0].islower()):
        df_twitter_arch_clean.loc[index, 'name_clean']='None'

In [0]: # TEST Issue 12
df_twitter_arch_clean[df_twitter_arch_clean['name_clean'].str[0].str.islower()]['name_cl

Out[0]: Series([], Name: name_clean, dtype: int64)

In [0]: df_twitter_arch_clean.name_clean.value_counts()

Out[0]: None          731
        Lucy           11
        Charlie        11
        Cooper         10
        Oliver         10

```

Tucker	9
Penny	9
Sadie	8
Winston	8
Lola	8
Daisy	7
Toby	7
Bella	6
Bo	6
Jax	6
Bailey	6
Koda	6
Stanley	6
Oscar	6
Bentley	5
Buddy	5
Louis	5
Rusty	5
Milo	5
Leo	5
Chester	5
Dave	5
Scout	5
Brody	4
Oakley	4
...	
Rooney	1
Jed	1
Dotsy	1
Ester	1
Wishes	1
Patch	1
Raphael	1
Yukon	1
Kendall	1
Baloo	1
Alf	1
Nugget	1
Robin	1
Zuzu	1
Berkeley	1
Rizzy	1
Pipsy	1
Lillie	1
Venti	1
Zeek	1
Marvin	1
Zoe	1

```

Bayley      1
Iggy        1
Willem      1
Mo          1
Ivar        1
Leela       1
Sprinkles   1
Jebberston  1
Name: name_clean, Length: 931, dtype: int64

```

```

In [0]: #DEFINE Issue 11 Quality Issue - The dog name column has 'None'. Should be Blank instead
# CLEAN Issue 11
df_twitter_arch_clean.name_clean.replace(['None'], '', inplace=True)

```

```

In [0]: #TEST Issue 11
df_twitter_arch_clean[df_twitter_arch_clean.name_clean=='None'].shape[0]

```

```

Out[0]: 0

```

```

In [0]: # DEFINE: 14 Tidiness Issue - The predicted p1, p2 and p3 have values which are not dogs
# Relevant data would be to pick the prediction which is actually a dog with the highest
# Create a scoring system to ascertain the best probability of a dog

```

```

# CLEAN: 14

```

```

truefalsemap = {
    True: 1,
    False: 0
}

```

```

df_image_pred_clean['p1_score'] = df_image_pred_clean['p1_conf'] * df_image_pred_clean['p1_prob']
df_image_pred_clean['p2_score'] = df_image_pred_clean['p2_conf'] * df_image_pred_clean['p2_prob']
df_image_pred_clean['p3_score'] = df_image_pred_clean['p3_conf'] * df_image_pred_clean['p3_prob']
df_image_pred_clean.head()

```

```

Out[0]:
   tweet_id  ...  p3_score
0  666020888022790149  ...  0.061428
1  666029285002620928  ...  0.072010
2  666033412701032449  ...  0.116197
3  666044226329800704  ...  0.222752
4  666049248165822465  ...  0.154629

```

```

[5 rows x 15 columns]

```

```

In [0]: # CLEAN: 14 Contd.

```

```

# Store the data in a separate column called predicted_dog
# in case none of the prediction is a dog, keep it blank
df_image_pred_clean['predicted_dog'] = ''

```

```

for index, row in df_image_pred_clean.iterrows():
    if((row['p1_score']>row['p2_score']) & (row['p1_score']>row['p3_score'])):
        df_image_pred_clean.loc[index, 'predicted_dog']=row['p1']
    elif((row['p2_score']>row['p1_score']) & (row['p2_score']>row['p3_score'])):
        df_image_pred_clean.loc[index, 'predicted_dog']=row['p2']
    elif((row['p3_score']>row['p1_score']) & (row['p3_score']>row['p3_score'])):
        df_image_pred_clean.loc[index, 'predicted_dog']=row['p3']

```

```

In [0]: # TEST 14
df_image_pred_clean['predicted_dog'].value_counts()

```

```

Out[0]:
golden_retriever      388
Labrador_retriever    166
Pembroke               94
Chihuahua              94
pug                    61
toy_poodle             48
chow                   47
Samoyed                45
Pomeranian             41
malamute               32
French_bulldog         32
cocker_spaniel         31
Chesapeake_Bay_retriever 30
miniature_pinscher     26
Cardigan               22
Staffordshire_bullterrier 22
German_shepherd        21
Eskimo_dog             20
Siberian_husky         20
Shih-Tzu               20
Rottweiler             19
Lakeland_terrier       18
beagle                 18
Shetland_sheepdog      18
kuvasz                 18
Maltese_dog            18
Italian_greyhound      17
basset                 17
West_Highland_white_terrier 16
...
giant_schnauzer        4
Saluki                 4
Gordon_setter          4
Ibizan_hound           4
Weimaraner             4
Afghan_hound           4

```

Rhodesian_ridgeback	4
Leonberg	3
cairn	3
komondor	3
Brabancon_griffon	3
toy_terrier	3
briard	3
Irish_water_spaniel	3
curly-coated_retriever	3
Scottish_deerhound	3
Greater_Swiss_Mountain_dog	3
Sussex_spaniel	2
Appenzeller	2
groenendael	2
Australian_terrier	2
black-and-tan_coonhound	2
wire-haired_fox_terrier	2
Japanese_spaniel	1
standard_schnauzer	1
silky_terrier	1
Irish_wolfhound	1
EntleBucher	1
clumber	1
Scotch_terrier	1

Name: predicted_dog, Length: 113, dtype: int64

In [0]: *# DEFINE: 16 Tidiness Issue - Pull in the data from Image Prediction and Json data to the*

CLEAN 16

Lets put all the data together in the twitter archive dataframe

Merge Archived tweet data with the downloaded Json data

Note, here we will want only the tweets which are not re-tweets which is what we have

Therefore use and innerjoin where entries from both twitter archive and json data match

`df_master_data = pd.merge(df_twitter_arch_clean, df_jdata_clean, on='tweet_id', how='inner')`

For the image prediction merge, we will do a left join to get the details for the main

`df_master_data = pd.merge(df_master_data, df_image_pred_clean, on='tweet_id', how='left')`

In [0]: *# TEST 16*

`df_master_data.info()`

<class 'pandas.core.frame.DataFrame'>

Int64Index: 2121 entries, 0 to 2120

Data columns (total 41 columns):

tweet_id	2121 non-null int64
in_reply_to_status_id	77 non-null float64
in_reply_to_user_id	77 non-null float64

```

timestamp                2121 non-null object
source                   2121 non-null object
text                     2121 non-null object
retweeted_status_id      0 non-null float64
retweeted_status_user_id 0 non-null float64
retweeted_status_timestamp 0 non-null object
expanded_urls            2063 non-null object
rating_numerator         2121 non-null int64
rating_denominator       2121 non-null int64
name                     2121 non-null object
doggo                    2121 non-null object
floofer                  2121 non-null object
pupper                   2121 non-null object
puppo                    2121 non-null object
combined_puppy_stage     2121 non-null object
text_uppercase           2121 non-null object
numerator_clean          2121 non-null float64
denominator_clean        2121 non-null float64
timestamp_clean          2121 non-null datetime64[ns, UTC]
retweeted_status_timestamp_clean 0 non-null datetime64[ns]
name_clean               2121 non-null object
retweet_count            2121 non-null int64
favorite_count           2121 non-null int64
jpg_url                  1940 non-null object
img_num                  1940 non-null float64
p1                       1940 non-null object
p1_conf                  1940 non-null float64
p1_dog                   1940 non-null object
p2                       1940 non-null object
p2_conf                  1940 non-null float64
p2_dog                   1940 non-null object
p3                       1940 non-null object
p3_conf                  1940 non-null float64
p3_dog                   1940 non-null object
p1_score                 1940 non-null float64
p2_score                 1940 non-null float64
p3_score                 1940 non-null float64
predicted_dog            1940 non-null object
dtypes: datetime64[ns, UTC](1), datetime64[ns](1), float64(13), int64(5), object(21)
memory usage: 696.0+ KB

```

```

In [0]: # DEFINE: 10. Tidiness Issue Lets drop some of the additional columns to make it look tidy

# CLEAN
cols_to_drop = ['in_reply_to_status_id', 'in_reply_to_user_id', 'doggo', 'floofer', 'pupper', 'puppo',
               'combined_puppy_stage', 'text_uppercase', 'numerator_clean', 'denominator_clean',
               'timestamp_clean', 'retweeted_status_timestamp_clean', 'name_clean', 'retweet_count',
               'favorite_count', 'jpg_url', 'img_num', 'p1', 'p1_conf', 'p1_dog', 'p2', 'p2_conf',
               'p2_dog', 'p3', 'p3_conf', 'p3_dog', 'p1_score', 'p2_score', 'p3_score', 'predicted_dog']
df_mastr_clean = df_master_data.copy()
df_mastr_clean.drop(cols_to_drop, axis=1, inplace=True)

```

```
In [0]: # TEST
df_mastr_clean.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 2121 entries, 0 to 2120
Data columns (total 21 columns):
tweet_id          2121 non-null int64
timestamp         2121 non-null object
source           2121 non-null object
text             2121 non-null object
rating_numerator  2121 non-null int64
rating_denominator 2121 non-null int64
name             2121 non-null object
combined_puppy_stage 2121 non-null object
text_uppercase   2121 non-null object
numerator_clean  2121 non-null float64
denominator_clean 2121 non-null float64
timestamp_clean   2121 non-null datetime64[ns, UTC]
name_clean       2121 non-null object
retweet_count    2121 non-null int64
favorite_count   2121 non-null int64
jpg_url          1940 non-null object
img_num          1940 non-null float64
p1_score         1940 non-null float64
p2_score         1940 non-null float64
p3_score         1940 non-null float64
predicted_dog    1940 non-null object
dtypes: datetime64[ns, UTC](1), float64(6), int64(5), object(9)
memory usage: 364.5+ KB
```

```
In [0]: df_mastr_clean.head()
```

```
Out[0]:
```

	tweet_id	timestamp	...	p3_score	predicted_dog
0	892420643555336193	2017-08-01 16:23:56 +0000	...	0.000000	
1	892177421306343426	2017-08-01 00:17:27 +0000	...	0.068957	Chihuahua
2	891815181378084864	2017-07-31 00:18:03 +0000	...	0.031379	Chihuahua
3	891689557279858688	2017-07-30 15:58:51 +0000	...	0.000000	Labrador_retriever
4	891327558926688256	2017-07-29 16:00:24 +0000	...	0.175219	basset

[5 rows x 21 columns]

5 Analyze and visualize - three (3) insights and one (1) visualization

```
In [0]: # What is the average dog rating
df_mastr_clean.numerator_clean.mean()
```

```
Out[0]: 12.161848184818481
```

```

In [0]: # What is the average rating by dog type as based on the image predictions
df_mastr_clean[df_mastr_clean.predicted_dog!=''].groupby('predicted_dog')['numerator_clean'].mean()
# Insight 1: This shows that the dog type Clumber most probably has the highest average rating

Out[0]: predicted_dog
Saluki      12.5
Name: numerator_clean, dtype: float64

In [0]: df_mastr_clean[df_mastr_clean.predicted_dog!=''].groupby('predicted_dog')['numerator_clean'].min()
# Insight 1 contd. the dog type Japanese Spaniel most probably has the lowest average rating

Out[0]: predicted_dog
Japanese_spaniel    5.0
Name: numerator_clean, dtype: float64

In [0]: # Insight 2: Which type of dogs were retweeted the most

df_mastr_clean[df_mastr_clean.predicted_dog!=''].groupby('predicted_dog')['retweet_count'].max()

# Bedlington Terrier dogs were the most re-tweeted

Out[0]: predicted_dog
Bedlington_terrier    8740.2
Name: retweet_count, dtype: float64

In [0]: # Insight 2: Contd.: Which type of dogs were retweeted the least

df_mastr_clean[df_mastr_clean.predicted_dog!=''].groupby('predicted_dog')['retweet_count'].min()

# groenendael dogs were the least re-tweeted

Out[0]: predicted_dog
groenendael    276.5
Name: retweet_count, dtype: float64

In [0]: # Insight 3: Which type of dogs were favorited the most

df_mastr_clean.groupby('predicted_dog')['favorite_count'].mean().sort_values(ascending=False)

# Saluki dogs were the most favorited

Out[0]: predicted_dog
Bedlington_terrier    24438.4
Name: favorite_count, dtype: float64

In [0]: # Insight 3: Contd.: Which type of dogs were favorited the least

df_mastr_clean.groupby('predicted_dog')['favorite_count'].mean().sort_values(ascending=True)

# Brabancon-griffon dogs were the least favorited

Out[0]: predicted_dog
Brabancon_griffon    885.0
Name: favorite_count, dtype: float64

```



```

In [0]: #Insight 4: Which is the most common dog name
df_mastr_clean['name_clean'].value_counts().sort_values(ascending=False).head(5)

# Charlie and Lucy seems to be the most common dog names followed by Cooper and Oliver

Out[0]:
Charlie      11
Lucy         11
Oliver       10
Cooper       10
Name: name_clean, dtype: int64

In [0]: #Insight 5: which is the most favorited dog stage
df_mastr_clean.groupby('combined_puppy_stage')['favorite_count'].mean().sort_values(ascending=False)

#puppo seems to be most popular

Out[0]: combined_puppy_stage
doggo,puppo    47844.0
Name: favorite_count, dtype: float64

In [0]: # When was the first tweet and the last tweet in this dataset
df_mastr_clean.timestamp_clean.min()

Out[0]: Timestamp('2015-11-15 22:32:08+0000', tz='UTC')


In [0]: df_mastr_clean.timestamp_clean.max()

Out[0]: Timestamp('2017-08-01 16:23:56+0000', tz='UTC')

In [145]: #Visualization 1
# What was the frequency of tweets based on this dataset
ax = df_mastr_clean.groupby(df_mastr_clean['timestamp_clean'].dt.strftime('%Y %b'))['text'].count()
ax.set_title('Tweeting Frequency')
ax.set_xlabel('Timeline')
ax.set_ylabel('No. of Tweets')

Out[145]: Text(0, 0.5, 'No. of Tweets')

```




```
In [146]: # Visualization 2
          # Which type of dog was the most tweeted about. Since there are many dog types
          # Lets pare it down to the top 10 most frequently tweeted dog type
df_tweet_countby_dogtype = df_mastr_clean[df_mastr_clean['predicted_dog'] != ''].groupby(
df_tweet_countby_dogtype
```

```
Out[146]: predicted_dog
golden_retriever      150
Labrador_retriever    99
Pembroke              90
Chihuahua             89
pug                   55
toy_poodle            47
Pomeranian            41
chow                  41
Samoyed               37
malamute              30
Name: tweet_id, dtype: int64
```

```
In [144]: # lets do a pie chart for this information
ax = df_tweet_countby_dogtype.plot.pie(autopct='%1.0f%%')
ax.set_title('Percentage of tweets by Dog Breed')
```

```
Out[144]: Text(0.5, 1.0, 'Percentage of tweets by Dog Breed')
```



```
In [0]: df_mastr_clean.to_csv(fn_final_clean_file, index=False)
```

```
In [0]:
```